

# Project\_2\_CemKazan

February 28, 2023

## 1 Structural Discovery of Macromolecules

author: Cem Kazan date: 02/28/2023

## 2 Structural Discovery of Macromolecules

The main question in structural discovery field is how macro molecules behave in their natural forms and natural environments and how macromolecules stray away from natural behavior.

## 3 Setup for Data Analysis

## Loading Libraries and the Data Frame

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from sklearn.cluster import DBSCAN
from sklearn.cluster import KMeans
from sklearn.cluster import AgglomerativeClustering
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
sns.set_context('talk')

[2]: import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from keras.wrappers.scikit_learn import KerasClassifier
```

- For data visualization, I use **matplotlib** or **seaborn** to uncover patterns and relationships in my data.
- For data manipulation, I use **pandas** to filter, summarize, and transform my data.
- For clustering analysis, I use **scikit-learn** and its **DBSCAN** algorithm to identify clusters within large datasets.
- For deep learning, I use **TensorFlow** to build and train neural networks
- For dimensionality reduction, I use **scikit-learn** and its **PCA** algorithm to reduce the number of features in my data while retaining as much information as possible

[2]:

### 3.1 Data Summary

```
[3]: df_input = pd.read_csv('./pdb_data_no_dups.csv')
display(df_input)
df_input.describe(include='all')
```

	structureId	classification	\
0	100D	DNA-RNA HYBRID	
1	101D	DNA	
2	101M	OXYGEN TRANSPORT	
3	102D	DNA	
4	102L	HYDROLASE(O-GLYCOSYL)	
...	...	...	
141396	9RUB	LYASE(CARBON-CARBON)	
141397	9TNA	T-RNA	
141398	9WGA	LECTIN (AGGLUTININ)	
141399	9XIA	ISOMERASE(INTRAMOLECULAR OXIDOREDUCTASE)	
141400	9XIM	ISOMERASE(INTRAMOLECULAR OXIDOREDUCTASE)	

	experimentalTechnique	macromoleculeType	residueCount	resolution	\
0	X-RAY DIFFRACTION	DNA/RNA Hybrid	20	1.90	
1	X-RAY DIFFRACTION	DNA	24	2.25	
2	X-RAY DIFFRACTION	Protein	154	2.07	
3	X-RAY DIFFRACTION	DNA	24	2.20	
4	X-RAY DIFFRACTION	Protein	165	1.74	
...	...	...	...	...	
141396	X-RAY DIFFRACTION	Protein	932	2.60	
141397	X-RAY DIFFRACTION	NaN	0	NaN	
141398	X-RAY DIFFRACTION	Protein	342	1.80	
141399	X-RAY DIFFRACTION	Protein	388	1.90	
141400	X-RAY DIFFRACTION	Protein	1572	2.40	

	structureMolecularWeight	crystallizationMethod	\
0	6360.30	VAPOR DIFFUSION, HANGING DROP	
1	7939.35	NaN	
2	18112.80	NaN	
3	7637.17	VAPOR DIFFUSION, SITTING DROP	
4	18926.61	NaN	
...	...	...	
141396	101838.68	NaN	
141397	24244.34	NaN	
141398	34270.22	NaN	
141399	43542.29	NaN	
141400	174722.12	NaN	

	crystallizationTempK	densityMatthews	densityPercentSol	\
0	NaN	1.78	30.89	
1	NaN	2.00	38.45	
2	NaN	3.09	60.20	
3	277.0	2.28	46.06	
4	NaN	2.75	55.28	
...	...	...	...	
141396	NaN	2.38	48.29	
141397	NaN	3.17	61.18	
141398	NaN	2.50	50.76	
141399	NaN	2.79	55.93	
141400	NaN	3.96	68.92	

	pdbxDetails	phValue	\
0	pH 7.00, VAPOR DIFFUSION, HANGING DROP	7.0	
1	NaN	NaN	
2	3.0 M AMMONIUM SULFATE, 20 MM TRIS, 1MM EDTA, ...	9.0	
3	pH 7.00, VAPOR DIFFUSION, SITTING DROP, temper...	7.0	
4	NaN	NaN	
...	...	...	
141396	NaN	NaN	
141397	NaN	NaN	
141398	NaN	NaN	
141399	NaN	NaN	
141400	NaN	NaN	

	publicationYear
0	1994.0
1	1995.0
2	1999.0
3	1995.0
4	1993.0
...	...
141396	1991.0
141397	1986.0
141398	1990.0
141399	1989.0
141400	1992.0

[141401 rows x 14 columns]

[3]:	structureId	classification	experimentalTechnique	macromoleculeType	\
count	141401	141399	141401	137636	
unique	140911	5050	33	13	
top	2FYM	HYDROLASE	X-RAY DIFFRACTION	Protein	
freq	4	20915	126432	127798	
mean	NaN	NaN	NaN	NaN	

std	NaN	NaN	NaN	NaN
min	NaN	NaN	NaN	NaN
25%	NaN	NaN	NaN	NaN
50%	NaN	NaN	NaN	NaN
75%	NaN	NaN	NaN	NaN
max	NaN	NaN	NaN	NaN

	residueCount	resolution	structureMolecularWeight	\
count	141401.000000	128589.000000	1.414010e+05	
unique	NaN	NaN	NaN	
top	NaN	NaN	NaN	
freq	NaN	NaN	NaN	
mean	825.374849	2.263807	1.120790e+05	
std	2136.461080	1.410878	5.690152e+05	
min	0.000000	0.480000	3.143800e+02	
25%	226.000000	1.800000	2.612856e+04	
50%	414.000000	2.100000	4.747779e+04	
75%	820.000000	2.500000	9.408484e+04	
max	313236.000000	70.000000	9.773054e+07	

	crystallizationMethod	crystallizationTempK	densityMatthews	\
count		96242	97039.000000	124724.000000
unique		549	NaN	NaN
top	VAPOR DIFFUSION, HANGING DROP		NaN	NaN
freq		53870	NaN	NaN
mean		NaN	290.967713	2.670267
std		NaN	9.541080	0.783740
min		NaN	4.000000	0.000000
25%		NaN	290.000000	2.210000
50%		NaN	293.000000	2.490000
75%		NaN	295.000000	2.910000
max		NaN	398.000000	99.000000

	densityPercentSol	pdboxDetails	phValue	publicationYear
count	124749.000000	118534	105110.000000	117602.000000
unique	NaN	91025	NaN	NaN
top	NaN	pH 7.5	NaN	NaN
freq	NaN	361	NaN	NaN
mean	51.353163	NaN	6.788685	2008.922365
std	10.104561	NaN	2.556819	8.459286
min	0.000000	NaN	0.000000	201.000000
25%	44.370000	NaN	6.000000	2005.000000
50%	50.500000	NaN	7.000000	2010.000000
75%	57.710000	NaN	7.500000	2014.000000
max	92.000000	NaN	724.000000	2018.000000

```
[4]: col_names = df_input.columns.tolist()
col_names_str = ', '.join(col_names)
```

The Structural Protein Sequences data [ @structur ] frame comprises 14 columns, as indicated by structureId, classification, experimentalTechnique, macromoleculeType, residueCount, resolution, structureMolecularWeight, crystallizationMethod, crystallizationTempK, densityMatthews, densityPercentSol, pdbxDetails, pHValue, publicationYear. An examination of the summary statistics revealed the presence of NA and empty values that must be addressed. Additionally, columns deemed irrelevant to our analysis have been removed.

```
[5]: df_clean = df_input.dropna()
df_clean = df_clean.drop(columns=['publicationYear', 'pdbxDetails'])
df_clean
```

```
[5]:
```

	structureId	classification	experimentalTechnique	macromoleculeType	\
3	102D	DNA	X-RAY DIFFRACTION	DNA	
27	110D	DNA	X-RAY DIFFRACTION	DNA	
30	111D	DNA	X-RAY DIFFRACTION	DNA	
36	113D	DNA	X-RAY DIFFRACTION	DNA	
44	117D	DNA	X-RAY DIFFRACTION	DNA	
...	...	...	...	...	
141064	6F6S	VIRAL PROTEIN	X-RAY DIFFRACTION	Protein	
141066	6F73	FLAVOPROTEIN	X-RAY DIFFRACTION	Protein	
141068	6F8P	VIRAL PROTEIN	X-RAY DIFFRACTION	Protein	
141077	6FAH	FLAVOPROTEIN	X-RAY DIFFRACTION	Protein	
141172	7BNA	DNA	X-RAY DIFFRACTION	DNA	

	residueCount	resolution	structureMolecularWeight	\
3	24	2.20	7637.17	
27	6	1.90	2337.73	
30	24	2.25	7374.83	
36	24	2.50	7356.81	
44	12	2.55	3663.39	
...	...	...	...	
141064	497	2.29	58337.03	
141066	1148	2.22	128053.03	
141068	316	1.60	34958.86	
141077	2074	3.13	231360.91	
141172	24	1.90	7326.78	

	crystallizationMethod	crystallizationTempK	densityMatthews	\
3	VAPOR DIFFUSION, SITTING DROP	277.00	2.28	
27	VAPOR DIFFUSION, SITTING DROP	277.00	2.90	
30	VAPOR DIFFUSION, SITTING DROP	277.00	2.29	
36	VAPOR DIFFUSION, SITTING DROP	281.00	2.35	
44	VAPOR DIFFUSION	277.00	3.01	
...	...	...	...	

141064	VAPOR DIFFUSION, SITTING DROP	293.00	3.83
141066	VAPOR DIFFUSION, HANGING DROP	294.00	2.90
141068	VAPOR DIFFUSION, SITTING DROP	298.00	2.47
141077	VAPOR DIFFUSION, SITTING DROP	293.15	3.49
141172	VAPOR DIFFUSION	290.00	2.27

	densityPercentSol	phValue
3	46.06	7.0
27	57.63	6.6
30	46.25	6.6
36	47.59	7.4
44	59.09	6.5
...	...	...
141064	67.89	5.2
141066	57.00	5.0
141068	50.15	7.0
141077	64.73	8.5
141172	45.79	7.5


[65886 rows x 12 columns]

## 4 Data Analysis


### 4.1 Molecule Type and Experiment Technique

First thing that need to be considered are the molecule types and the technique used for the experiment to choose ideal data for examination.

```
[6]: sns.countplot(y='macromoleculeType', data=df_clean)
plt.title("No of Data Points by Molecule Type")
plt.xlabel("Molecule Type")
plt.ylabel("Count")
plt.xticks(rotation=45)
plt.show()
```



```
[7]: sns.countplot(y='experimentalTechnique', data=df_clean)
plt.title("No of Data Points by Experimental Techniques")
plt.xlabel("Molecule Type")
plt.ylabel("Count")
plt.xticks(rotation=45)
plt.show()
```



```
[8]: df_select_1 = df_clean[
    (
        df_clean['macromoleculeType'] == 'Protein'
    ) & (
```

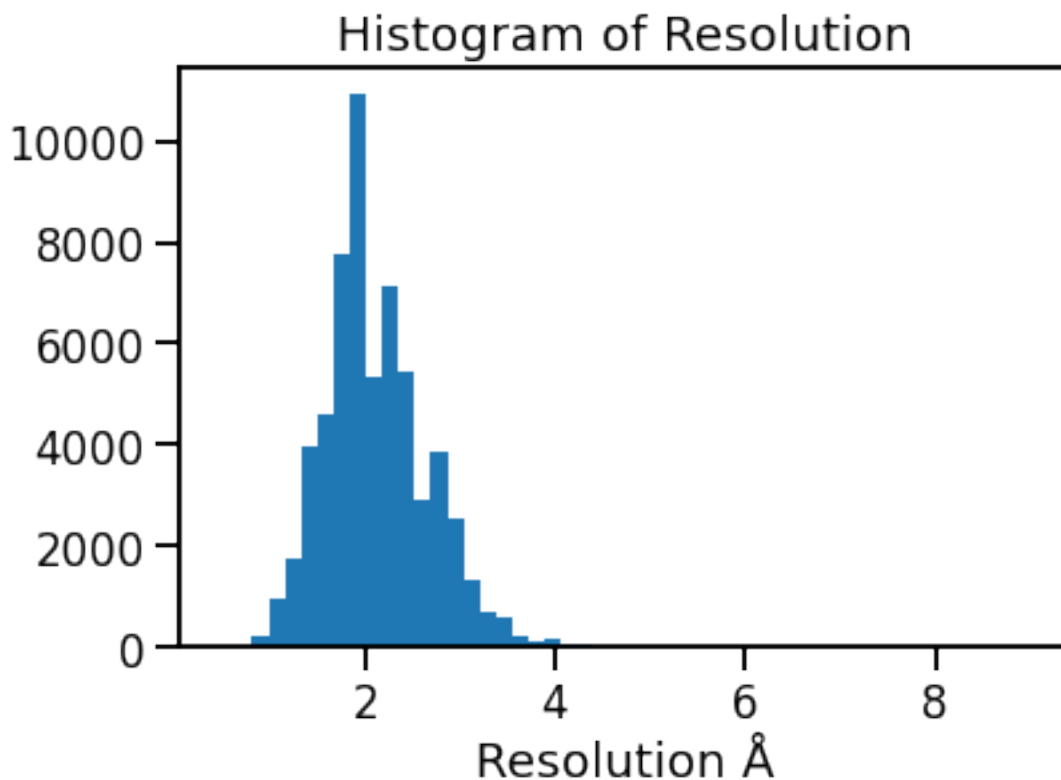
```
df_clean['experimentalTechnique'] == 'X-RAY DIFFRACTION'  
)  
].copy()
```

Due to the large amount of data available for proteins, it was selected as the subject of further analysis. Given the majority of the results are obtained through x-ray diffraction, my focus was specifically directed towards proteins that have been studied using this technique.

## 4.2 Effect of Resolution

Majority of the protein structural discoveries uses the x-ray diffraction and one major achievement in discovery is obtaining a resolution (in angstrom units) as close to one as possible.[@warren1990; @whittig1986] This is limited by the physics behind it. To understand this one dimensional data I used histogram to see how it is distributed.

```
[9]: plt.hist(df_select_1['resolution'], bins=50)  
plt.title("Histogram of Resolution")  
plt.xlabel("Resolution Å")  
plt.show()
```



The histogram of resolution shows that majority of the experiments can resolve structures with a resolution around 2. (As there are more than 60000 data points in my set, the number of bins is selected as 50 to get the details about the distribution more precisely). Investigation of structures



with resolution below 2 would give much deeper understanding of the structural features and the underlying physics/chemistry.

```
[10]: df_select_2 = df_select_1[
      (
          df_select_1['resolution'] < 2
      )
  ]
```

The number of protein structures with a resolution below two are 26248.

### 4.3 Effect of Residue Count



Next I wanted to understand how residue count (the number of residue in macromolecules) affects the resolution of the discovered structure.

```
[11]: df_protein1 = df_select_1[(df_select_1['resolution'] < 2) &
    ↪(df_select_1['residueCount'] < 2000)].copy()
df_protein2 = df_select_1[(df_select_1['resolution'] > 2) &
    ↪(df_select_1['residueCount'] < 2000)].copy()
df_protein1['label'] = 'resolution < 2'
df_protein2['label'] = 'resolution > 2'

df_plot = pd.concat([df_protein1, df_protein2])

sns.displot(
    data=df_plot,
    x='residueCount',
    hue='label',
    kind='hist',
    common_norm=False,
    binwidth=50,
    multiple='dodge',
    shrink=0.8,
    aspect=2,
)
sns.displot(
    data=df_plot,
    x='residueCount',
    hue='label',
    kind='kde',
    common_norm=False,
    aspect=2,
)
```

```
[11]: <seaborn.axisgrid.FacetGrid at 0x7f646b965f10>
```



The relation between resolution and number of residues is intriguing. The plot shows that number of residues in a protein is a factor in determining the protein resolution. There is a negative correlation between the number of residues and resolution. The lesser the number of residues better the resolution (close to 1). Most of the high resolution structures have less than 500 residues. Thus, I investigated that set.

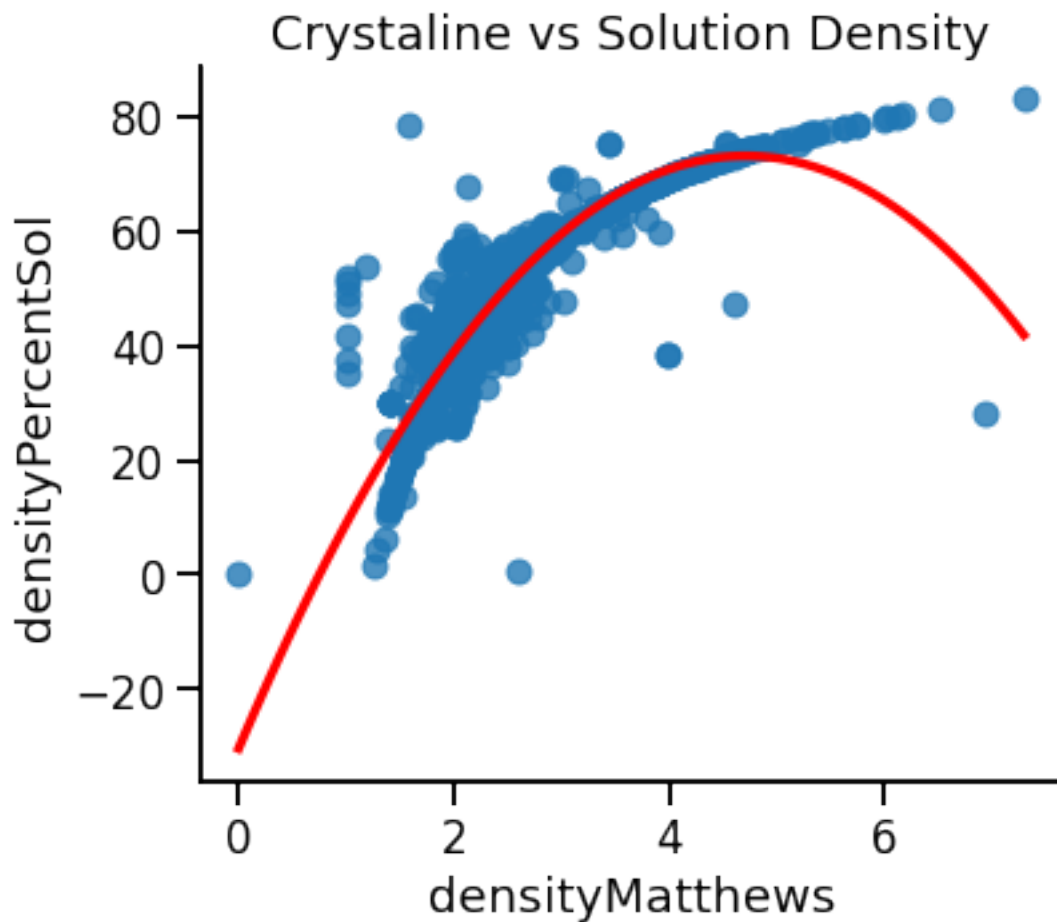
```
[12]: df_select_residu = df_select_2[(df_select_2['residueCount'] < 500)]
```

#### 4.4 Exploring Discrepancies in the Density of Proteins in Crystalline and Soluble Forms

Proteins in solution has shown to display features that are being missed in structural discovery. One of the parameters that has been trusted in protein studies in solution is solution density calculation. On the other side structural discovery utilizes a metric depending on the protein crystals. Now the question is how do proteins act in solution (normal conditions) versus in a crystal lattice (restricted environment) ? Thus I investigated two features from both calculations.

```
[13]: sns.lmplot(
    data = df_select_residu,
    x='densityMatthews',
    y='densityPercentSol',
    order = 2,
    ci = 0,
    line_kws={'color': 'red'},
    aspect=1.2,
    legend = True
)
plt.title("Crystalline vs Solution Density")
```

```
[13]: Text(0.5, 1.0, 'Crystalline vs Solution Density')
```



The spread of the data in the plot resembles a second degree polynomial fit, hence I added a second degree polynomial fit to capture the underlying principle that correlates these two metrics. However because of the outliers the fit does not recapitulate the data. Thus the outliers needs to be removed.

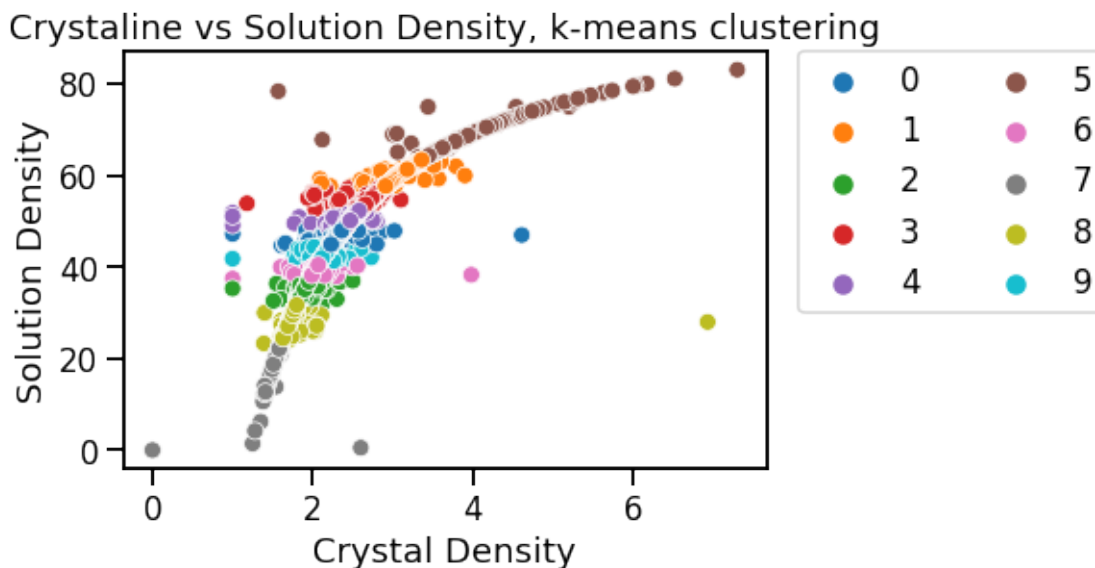
#### 4.4.1 K-means clustering

I used K-means clustering to discover similar data points in the clusters within the two feature space. It can reveal hidden patterns and trends such that it can cluster and separate the data points that don't follow the correct trend. There are thousands of data points in my system therefore I selected k value as 20 to be able to generate unique clusters.

```
[14]: df_density = df_select_residu[['densityMatthews', 'densityPercentSol']].copy()
kmeans_fit = KMeans(n_clusters=10, max_iter=20, n_init=2).fit(df_density)
df_density['clusters'] = kmeans_fit.labels_
```

```
[15]: sns.scatterplot(
    data = df_density,
    x='densityMatthews',
    y='densityPercentSol',
    legend = True,
    hue='clusters',
    palette = sns.color_palette('tab10')
)
plt.xlabel('Crystal Density')
plt.ylabel('Solution Density')
plt.title('Crystalline vs Solution Density, k-means clustering')
plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0., ncol=2)
```

```
[15]: <matplotlib.legend.Legend at 0x7f646b502be0>
```



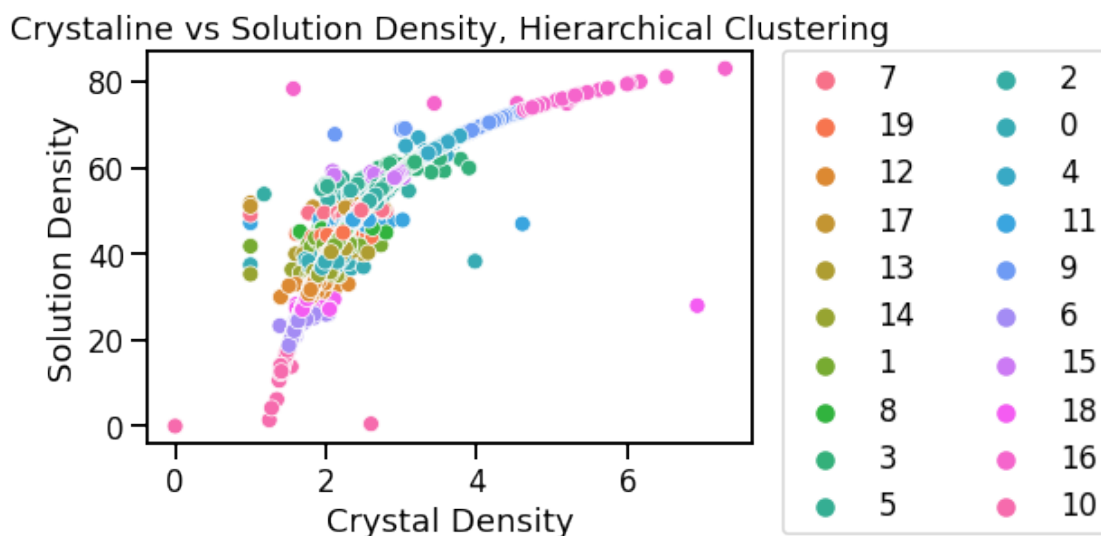
The K-means clustering failed to identify the correct clusters, therefore I used Hierarchical clustering that can identify clusters using agglomerative method which does not require a prior specification of the number of clusters.

#### 4.4.2 Hierarchical clustering

```
[16]: df_density_2 = df_select_residu[['densityMatthews', 'densityPercentSol']].copy()
hierarchical_cluster = AgglomerativeClustering(n_clusters=20,
        affinity='euclidean', linkage='ward')
labels = hierarchical_cluster.fit_predict(df_density_2)
df_density_2['cluster'] = labels
df_density_2['cluster'] = df_density_2['cluster'].astype(str)
```

```
[17]: sns.scatterplot(
    data = df_density_2,
    x='densityMatthews',
    y='densityPercentSol',
    legend = True,
    hue='cluster',
)
plt.xlabel('Crystal Density')
plt.ylabel('Solution Density')
plt.title('Crystalline vs Solution Density, Hierarchical Clustering')
plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0., ncol=2)
```

```
[17]: <matplotlib.legend.Legend at 0x7f646d452940>
```



The figures indicate that kmeans and hierarchical clustering was not the best approach to identify the main overlapping points versus the outliers in the data. This data show dense regions around several data points. Therefore a density based clustering approach would be more useful in understanding this data. To achieve that I used dbscan function. DBscan is a density-based clustering algorithm that groups together data points that are close in space and separates points that are far apart. Dbscan is useful because it can find clusters of arbitrary shapes and sizes, unlike k-means

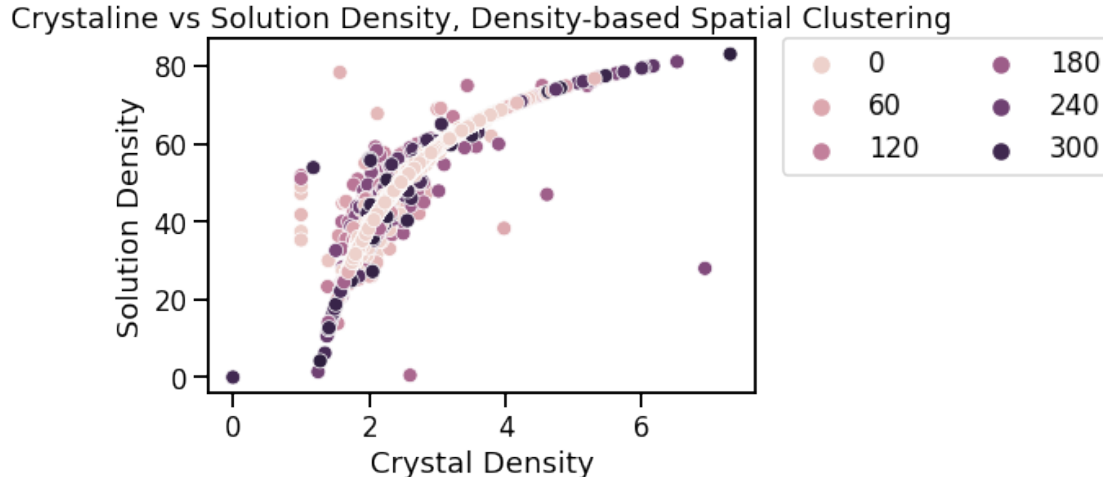
and hierarchical clustering which tend to find circular clusters.

#### 4.4.3 Density Based Clustering

```
[18]: df_density_3 = df_select_residu[['densityMatthews', 'densityPercentSol']].copy()
dbscan = DBSCAN(eps=0.1, min_samples=1)
result = dbscan.fit_predict(df_density_3)
df_density_3['cluster'] = result
# print(df_density_3['cluster'].value_counts())
```

```
[19]: sns.scatterplot(
    data = df_density_3,
    x='densityMatthews',
    y='densityPercentSol',
    legend = True,
    hue='cluster',
)
plt.xlabel('Crystal Density')
plt.ylabel('Solution Density')
plt.title('Crystalline vs Solution Density, Density-based Spatial Clustering')
plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0., ncol=2)
```

```
[19]: <matplotlib.legend.Legend at 0x7f646ab86b50>
```



Looking at the plot we can clearly see the outliers in our data set. To focus more clearly on the data I choose main cluster and re-plot the data. I added a second degree polynomial fit once again to discover the underlying relation. The polynomial fit now captures more than 95% of the data as shown in the plot below.

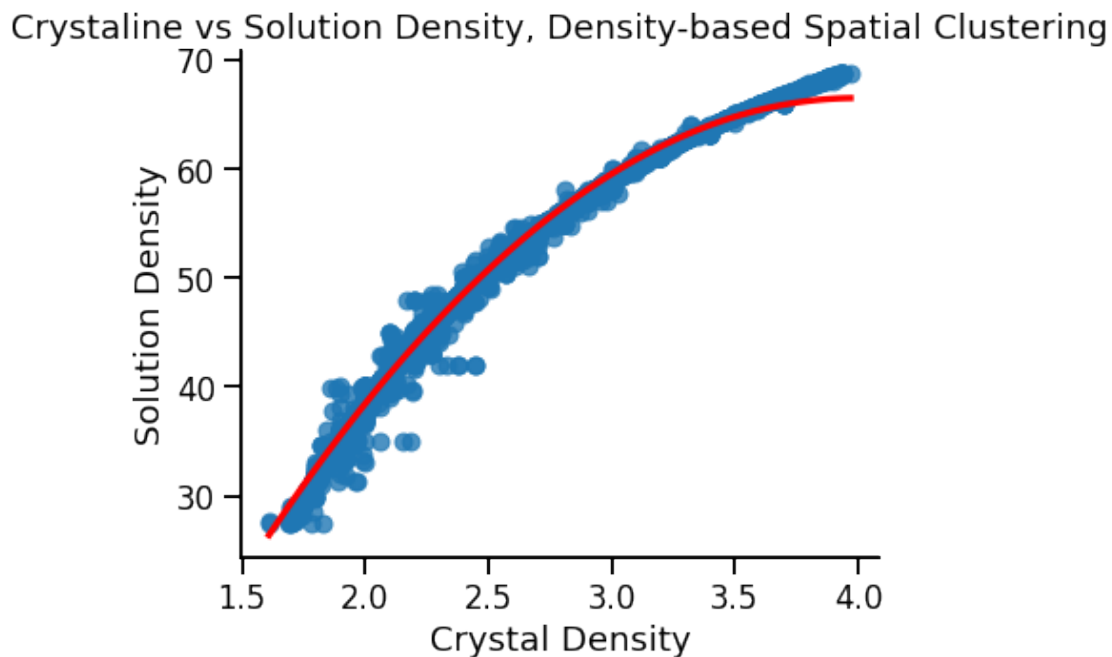
```
[20]: df_density_4 = df_density_3[df_density_3['cluster'] == 0]
```

```

sns.lmplot(
    data = df_density_4,
    x='densityMatthews',
    y='densityPercentSol',
    order = 2,
    ci = 0,
    line_kws={'color': 'red'},
    aspect=1.2,
)
plt.title("Crystalline vs Solution Density")
plt.xlabel('Crystal Density')
plt.ylabel('Solution Density')
plt.title('Crystalline vs Solution Density, Density-based Spatial Clustering')

```

[20]: Text(0.5, 1.0, 'Crystalline vs Solution Density, Density-based Spatial Clustering')



The results of this analysis demonstrate that through the utilization of various machine learning techniques, it is possible to identify and decipher fundamental relationships within a data set. This study highlights the challenges of achieving results that resemble natural phenomena using these techniques, with only 26% of the data points, equivalent to only 26% of the experiments, resulting in naturally occurring behavior that captures the correct underlying fundamentals.

So far focused on only several features from the dataset and filtering based on physics based

principals I was able to distinguish different protein structures from each other. However, the dataset contains vast number of features that could be understood in a cohesive manner. In example for every protein 5 features could be analyzed at the same time to achieve higher clarity in its natural behaviour relative with other proteins in the set. Several mathematical approaches could be utilized to achieve this. In the next chapters I will apply several different methods to analyze the data. This will hopefully provide a better understanding of the data and potentially be able to distinguish proteins with similar characteristics from the ones that behave differently.

## 5 Principal Component Analysis

PCA is a technique that helps us understand complex datasets by simplifying them into a smaller number of variables (dimension reduction). The most significant links and patterns in the data are found, and new variables are then created to capture those relationships and patterns. These new variables, referred to as principle components, are simpler to use and can provide us with new perspectives on the underlying structure of the data.

First I will define a python function that takes a dataframe with several features and outputs its principal components.

```
[21]: def pca_df(df):
    # first we need to scale the values to avoid bias
    x = StandardScaler(with_mean=True, with_std=False).fit_transform(df.values)
    # now we do pca model and fit
    pca = PCA()
    principal_components = pca.fit_transform(x)
    print(pca.singular_values_)
    principal_df = pd.DataFrame(
        data=principal_components,
        index=df.index,
        columns=['PC{}'.format(a) for a in range(1, principal_components.
↪shape[1] + 1)]
    )
    return principal_df
```

```
[22]: df_clean
```

```
[22]:      structureId classification experimentalTechnique macromoleculeType \
3           102D          DNA      X-RAY DIFFRACTION          DNA
27          110D          DNA      X-RAY DIFFRACTION          DNA
30          111D          DNA      X-RAY DIFFRACTION          DNA
36          113D          DNA      X-RAY DIFFRACTION          DNA
44          117D          DNA      X-RAY DIFFRACTION          DNA
...          ...          ...          ...          ...
141064      6F6S  VIRAL PROTEIN      X-RAY DIFFRACTION      Protein
141066      6F73  FLAVOPROTEIN      X-RAY DIFFRACTION      Protein
141068      6F8P  VIRAL PROTEIN      X-RAY DIFFRACTION      Protein
141077      6FAH  FLAVOPROTEIN      X-RAY DIFFRACTION      Protein
141172      7BNA          DNA      X-RAY DIFFRACTION          DNA
```



	residueCount	resolution	structureMolecularWeight	\
3	24	2.20	7637.17	
27	6	1.90	2337.73	
30	24	2.25	7374.83	
36	24	2.50	7356.81	
44	12	2.55	3663.39	
...	...	...	...	
141064	497	2.29	58337.03	
141066	1148	2.22	128053.03	
141068	316	1.60	34958.86	
141077	2074	3.13	231360.91	
141172	24	1.90	7326.78	

	crystallizationMethod	crystallizationTempK	densityMatthews	\
3	VAPOR DIFFUSION, SITTING DROP	277.00	2.28	
27	VAPOR DIFFUSION, SITTING DROP	277.00	2.90	
30	VAPOR DIFFUSION, SITTING DROP	277.00	2.29	
36	VAPOR DIFFUSION, SITTING DROP	281.00	2.35	
44	VAPOR DIFFUSION	277.00	3.01	
...	...	...	...	
141064	VAPOR DIFFUSION, SITTING DROP	293.00	3.83	
141066	VAPOR DIFFUSION, HANGING DROP	294.00	2.90	
141068	VAPOR DIFFUSION, SITTING DROP	298.00	2.47	
141077	VAPOR DIFFUSION, SITTING DROP	293.15	3.49	
141172	VAPOR DIFFUSION	290.00	2.27	

	densityPercentSol	phValue
3	46.06	7.0
27	57.63	6.6
30	46.25	6.6
36	47.59	7.4
44	59.09	6.5
...	...	...
141064	67.89	5.2
141066	57.00	5.0
141068	50.15	7.0
141077	64.73	8.5
141172	45.79	7.5

[65886 rows x 12 columns]

```
[23]: df_PCA_input= df_clean[[
        'structureId',
        'residueCount','resolution','structureMolecularWeight',
        'crystallizationTempK','densityMatthews','densityPercentSol','phValue'
    ]]
```

```
df_PCA_input.set_index('structureId', inplace=True)
display(df_PCA_input)
```

structureId	residueCount	resolution	structureMolecularWeight	\
102D	24	2.20	7637.17	
110D	6	1.90	2337.73	
111D	24	2.25	7374.83	
113D	24	2.50	7356.81	
117D	12	2.55	3663.39	
...	...	...	...	
6F6S	497	2.29	58337.03	
6F73	1148	2.22	128053.03	
6F8P	316	1.60	34958.86	
6FAH	2074	3.13	231360.91	
7BNA	24	1.90	7326.78	

structureId	crystallizationTempK	densityMatthews	densityPercentSol	phValue
102D	277.00	2.28	46.06	7.0
110D	277.00	2.90	57.63	6.6
111D	277.00	2.29	46.25	6.6
113D	281.00	2.35	47.59	7.4
117D	277.00	3.01	59.09	6.5
...	...	...	...	...
6F6S	293.00	3.83	67.89	5.2
6F73	294.00	2.90	57.00	5.0
6F8P	298.00	2.47	50.15	7.0
6FAH	293.15	3.49	64.73	8.5
7BNA	290.00	2.27	45.79	7.5

[65886 rows x 7 columns]

```
[24]: principal_df = pca_df(df_PCA_input)
print(principal_df)

df_PCA_output = pd.concat([df_PCA_input, principal_df], axis=1)
print(df_PCA_output)
```

```
[7.41476344e+07 1.63878371e+05 2.56273553e+03 2.29696160e+03
 3.31008249e+02 1.27457167e+02 6.82179473e+01]
      PC1      PC2      PC3      PC4      PC5 \
structureId
102D    -92952.188408  302.324139  1.442777  14.778515  0.219619
110D    -98251.650145  292.661844  12.073082  10.172944 -0.197828
111D    -93214.524838  300.955409  1.617215  14.702583 -0.179696
113D    -93232.544587  300.858965  1.262695  10.496945  0.626207
117D    -96925.976875  293.574242  13.425705   9.584843 -0.288906
```

...	...	...	...	...	...
6F6S	-42250.550792	93.752619	14.588273	-8.368569	-1.597463
6F73	27467.896436	-193.569594	3.390525	-4.616139	-1.794568
6F8P	-65629.346967	152.835223	-3.529095	-5.956155	0.218748
6FAH	130779.201212	-580.698352	9.772903	-6.473497	1.696199
7BNA	-93262.574191	300.705058	-3.983118	2.962916	0.725861

	PC6	PC7
structureId		
102D	0.270970	-0.071607
110D	-0.330770	-0.149403
111D	0.322020	-0.079975
113D	0.513798	-0.125839
117D	0.280719	-0.198232
...	...	...
6F6S	-0.226797	0.102477
6F73	-0.094823	-0.120307
6F8P	-0.499678	-0.064418
6FAH	0.482994	-0.080599
7BNA	-0.049173	-0.028479

[65886 rows x 7 columns]

	residueCount	resolution	structureMolecularWeight	\
structureId				
102D	24	2.20	7637.17	
110D	6	1.90	2337.73	
111D	24	2.25	7374.83	
113D	24	2.50	7356.81	
117D	12	2.55	3663.39	
...	...	...	...	
6F6S	497	2.29	58337.03	
6F73	1148	2.22	128053.03	
6F8P	316	1.60	34958.86	
6FAH	2074	3.13	231360.91	
7BNA	24	1.90	7326.78	

	crystallizationTempK	densityMatthews	densityPercentSol	\
structureId				
102D	277.00	2.28	46.06	
110D	277.00	2.90	57.63	
111D	277.00	2.29	46.25	
113D	281.00	2.35	47.59	
117D	277.00	3.01	59.09	
...	...	...	...	
6F6S	293.00	3.83	67.89	
6F73	294.00	2.90	57.00	
6F8P	298.00	2.47	50.15	
6FAH	293.15	3.49	64.73	

7BNA		290.00	2.27	45.79	
------	--	--------	------	-------	--

	pHValue	PC1	PC2	PC3	PC4 \
structureId					
102D	7.0	-92952.188408	302.324139	1.442777	14.778515
110D	6.6	-98251.650145	292.661844	12.073082	10.172944
111D	6.6	-93214.524838	300.955409	1.617215	14.702583
113D	7.4	-93232.544587	300.858965	1.262695	10.496945
117D	6.5	-96925.976875	293.574242	13.425705	9.584843
...	...	...	...	...	...
6F6S	5.2	-42250.550792	93.752619	14.588273	-8.368569
6F73	5.0	27467.896436	-193.569594	3.390525	-4.616139
6F8P	7.0	-65629.346967	152.835223	-3.529095	-5.956155
6FAH	8.5	130779.201212	-580.698352	9.772903	-6.473497
7BNA	7.5	-93262.574191	300.705058	-3.983118	2.962916


  

	PC5	PC6	PC7
structureId			
102D	0.219619	0.270970	-0.071607
110D	-0.197828	-0.330770	-0.149403
111D	-0.179696	0.322020	-0.079975
113D	0.626207	0.513798	-0.125839
117D	-0.288906	0.280719	-0.198232
...	...	...	...
6F6S	-1.597463	-0.226797	0.102477
6F73	-1.794568	-0.094823	-0.120307
6F8P	0.218748	-0.499678	-0.064418
6FAH	1.696199	0.482994	-0.080599
7BNA	0.725861	-0.049173	-0.028479

[65886 rows x 14 columns]

The first two principal components in PCA are often the most important because they capture the largest amount of variation in the original dataset. This means that they contain the most information about the underlying patterns and relationships in the data. By analyzing these components, we can gain a better understanding of the key factors that are driving the variation in the dataset, and use this information to make better decisions or predictions. Essentially, the first two components give us a high-level overview of the data and provide a foundation for deeper analysis.

```
[25]: plt.scatter(
        x= df_PCA_output['PC1'],
        y= df_PCA_output['PC2'],
    )
plt.axhline(y=0, color='k', linestyle='--', alpha=0.2)
plt.axvline(x=0, color='k', linestyle='--', alpha=0.2)
plt.xlabel('PC1')
plt.ylabel('PC2')
plt.show()
```



By plotting the first two principal components using matplotlib scatter plot we can identify data points that are most deviant and the ones that are similar to each other. Here intriguingly PC2 identifies and separates outliers very clearly. Thus first I will investigate the differences between these clusters.

```
[26]: df_PC2_1 = df_PCA_output [ (df_PCA_output['PC2'] > 0)]
      df_PC2_2 = df_PCA_output [ (df_PCA_output['PC2'] < 0)]
```

```
[27]: df_PC2_1.mean()
```

```
[27]: residueCount      435.409737
      resolution        2.030232
      structureMolecularWeight  60391.695490
      crystallizationTempK    291.161979
      densityMatthews      2.604698
      densityPercentSol    50.299365
      pHValue             6.773873
      PC1                 -40196.234610
      PC2                 166.088967
      PC3                 -0.758308
      PC4                  0.285095
      PC5                 -0.008981
      PC6                 -0.056183
      PC7                  0.013720
      dtype: float64
```

```
[28]: df_PC2_2.mean()
```

```
[28]: residueCount      1637.439114
      resolution        2.429300
      structureMolecularWeight  188973.750474
      crystallizationTempK      291.013852
      densityMatthews        2.836193
      densityPercentSol      54.425100
      pHValue              6.818693
      PC1                  88390.341108
      PC2                 -365.224768
      PC3                   1.667497
      PC4                 -0.626916
      PC5                   0.019749
      PC6                   0.123545
      PC7                 -0.030170
      dtype: float64
```

By looking at the mean values of each feature in two clusters separated by PC2, we observe a marginal difference in the number of residues of a protein (residueCount). The first cluster have an average residue count of 435, while the second cluster has an average of 1637. Following, the molecular weights of the protein (structureMolecularWeight) shows a similar trend. One of the most important components in structural discovery , resolution , shows distinction in these two clusters. We have previously seen the effect of residue count on resolution in this study and PCA was able to identify this difference robustly. The first cluster with an average less number of residues shows an average resolution of 2 angstrom, while the second cluster containing proteins with a large number of residues has an average of 2.4 angstrom. On the other hand the second principal component was not able to distinguish the difference of densities of proteins.

```
[29]: df_PC1_1 = df_PCA_output [ (df_PCA_output['PC1'] > 0)]
      df_PC1_2 = df_PCA_output [ (df_PCA_output['PC1'] < 0)]
      print(df_PC1_1.mean())
```

```
residueCount      2133.092425
resolution        2.520103
structureMolecularWeight  276029.527503
crystallizationTempK      291.091253
densityMatthews        2.904525
densityPercentSol      55.344906
pHValue           6.818821
PC1                175447.519219
PC2               -406.775345
PC3                1.998317
PC4               -0.883510
PC5                0.010809
PC6                0.154124
PC7               -0.024116
dtype: float64
```

```
[30]: print(df_PC1_2.mean())
```

```
residueCount      399.138125
resolution         2.041178
structureMolecularWeight  45903.344035
crystallizationTempK    291.123286
densityMatthews     2.606167
densityPercentSol    50.418430
pHValue            6.778241
PC1                -54684.578158
PC2                126.786279
PC3                -0.622848
PC4                 0.275378
PC5                -0.003369
PC6                -0.048038
PC7                 0.007517
dtype: float64
```

By investigating the first principal component we again see a similar trend as PC2. Thus I looked into the third principal component.

```
[31]: df_PC3_1 = df_PCA_output [ (df_PCA_output['PC3'] > 0)]
df_PC3_2 = df_PCA_output [ (df_PCA_output['PC3'] < 0)]
print(df_PC3_1.mean())
```

```
residueCount      848.059225
resolution         2.357352
structureMolecularWeight  103099.207839
crystallizationTempK    288.626577
densityMatthews     3.170135
densityPercentSol    59.624731
pHValue            6.791331
PC1                2512.849230
PC2                -23.800663
PC3                 8.330472
PC4                -0.898369
PC5                -0.004923
PC6                -0.007954
PC7                -0.032932
dtype: float64
```

```
[32]: print(df_PC3_2.mean())
```

```
residueCount      779.146510
resolution         1.979356
structureMolecularWeight  98405.916627
crystallizationTempK    293.275799
densityMatthews     2.249158
```

densityPercentSol	44.615448
phValue	6.784893
PC1	-2180.737672
PC2	20.655040
PC3	-7.229472
PC4	0.779636
PC5	0.004272
PC6	0.006902
PC7	0.028579

dtype: float64

The average of the features when separating the data by its third principal component shows that the density of a protein in solution (densityPercentSol) is different in two clusters. Following, the density of a protein in crystal (densityMatthews) shows a similar trend. We have previously shown the correlation between the densities in solution and crystal, and we are observing the same distinction with PC3 clustering of the data. The crystal and solution densities for the first and second cluster of PC3 is 3.1, 59.6, and 2.2, 44.6, respectively.

By utilizing PCA, we observe that the features residue count, resolution and molecular weight are tied together within the second and first principal component, while density related features are tied to the third principal component. This implies that one can harness the power of dimensionality reduction techniques to identify the physical characteristics of proteins and distinguish, and potentially predict behaviours of new proteins by looking into several of their features.

## 6 Non-Linear Data Analyses with Artificial Neural Network Treatment

Previous analyses have shown great promise in deciphering the structural data of proteins. However, using all the features of the data together most of the methods uses functions in a linear fashion and this challenges and restricts the possible highly accurate understanding and predictions of the data. Hence, I wanted to follow a more non linear approach which are Artificial Neural Networks. In this study I want to explore the data and try to pin down the underlying physics to be able to classify proteins by their structural data. One interesting way to achieve this is using unsupervised learning. Unsupervised learning is useful when we do not already have the defining labels for the data and trying to learn the representations of the data which will provide us with the classification of the proteins in this study. Autoencoders, which are a type of neural networks, are often used for unsupervised learning tasks, meaning that the data used to train the autoencoder does not have any labels. Instead, the autoencoder tries to learn the underlying structure of the data without any guidance from an outside source or predefined criteria. Once the autoencoder is trained, it can be used for various classifications of proteins. It is composed of two parts: an encoder and a decoder. The encoder takes the input data and compresses it into a lower-dimensional representation, and the decoder takes this compressed representation and reconstructs the data as closely to the original as possible. This systematic algorithm learns the most important parts of the data and uses it to reconstruct it. Here I apply a simple autoencoder to understand the structural data from structural discovery.

```
[33]: # this function separates data as features and labels.
      # This is not required in the current case as it is unsupervised but
```



```

# the function splitting the data requires labels, hence
# I used dummy data as labels to use it later on splitting.
def prepare_xy(dataset, x_label_list, y_label_list):
    x = dataset.loc[:, x_label_list]
    y = dataset.loc[:, y_label_list]
    print('x: {}'.format(x.shape))
    print('y: {}'.format(y.shape))
    return x, y

```

```

[34]: # the data going into a neural network needs to be renormalized for it to be
# functioning better, meaning a data distributed evenly between negative and
# positive values with a mean around zero, would make multiplying and summing
# up the weights and values much easier, leading to shorter convergence of
# neural network training and easier learning.
def scaler_1(x):
    from sklearn.preprocessing import StandardScaler
    sc = StandardScaler()
    sc.fit(x)
    xs = sc.transform(x)
    return sc, xs

```

```

[35]: # This function is used to split the data into training and testing.
# This is an important step as I am training the model with part of my data,
# and testing the ability of my model in capturing correct features with
# previously unseen data for the neural network.
# I am using 90% training, 10% testing splitting.
# 85% to 15% is also a common way to split, but as I have thousands of data
# points, I would already have enough data points to test the neural network
# with.
def splitter_1(x, y, test_size=0.1):
    from sklearn.model_selection import train_test_split
    display(y.apply(pd.value_counts))
    x_train, x_test, y_train, y_test = train_test_split(x, y, test_size =
    test_size)
    print('x_train: {}'.format(x_train.shape))
    display(x_train)
    print('x_test: {}'.format(x_test.shape))
    display(x_test)
    print('y_train: {}'.format(y_train.shape))
    display(y_train)
    display(y_train.apply(pd.value_counts))
    print('y_test: {}'.format(y_test.shape))
    display(y_test)
    display(y_test.apply(pd.value_counts))
    return x_train, x_test, y_train, y_test

```

```

[36]: # Here I build the architecture of my neural net
def build_keras_model(input_dim, output_dim):
    classifier = tf.keras.Sequential()
    # Input layer takes the input data and passes it to hidden layers
    classifier.add(layers.Dense(
        units=(0.6 * input_dim),
        kernel_initializer=tf.initializers.he_uniform(),
        activation=tf.keras.activations.relu,
        input_dim = input_dim
    ))
    classifier.add(layers.Dense(
        units=(0.2 * input_dim),
        kernel_initializer=tf.initializers.he_uniform(),
        activation=tf.keras.activations.relu
    ))
    # The number of nodes in each layer decreases initially and then increases.
    # this is to reduce the dimensions as much as possible and try to
    # recapitulate the inital data values later
    classifier.add(layers.Dense(
        units=(0.6 * input_dim),
        kernel_initializer=tf.initializers.he_uniform(),
        activation=tf.keras.activations.relu
    ))
    # The final layer is the output layer which has a sigmoid function.
    # Sigmoid is important here as it give a distribution of values between 0 and 1,
    # and these values can be translated back into real values later.
    classifier.add(layers.Dense(
        units=output_dim,
        kernel_initializer=tf.initializers.glorot_uniform(),
        activation=tf.keras.activations.sigmoid
    ))
    # the last part is the seelction of minimizer and loss function. Adam is a
    # common minimizer and MSE is a robust function to check how much the system is
    # closer to the original data
    classifier.compile(
        optimizer=tf.keras.optimizers.Adam(),
        loss=tf.keras.losses.MeanSquaredError(),
        metrics=[
            tf.keras.metrics.Accuracy(),
            tf.keras.metrics.BinaryAccuracy(),
            tf.keras.metrics.AUC(),
            tf.keras.metrics.TruePositives(),
            tf.keras.metrics.TrueNegatives(),

```

```

        tf.keras.metrics.FalsePositives(),
        tf.keras.metrics.FalseNegatives()
    ]
)
return classifier

```

```

[37]: # this function is to do the fitting of the data on to the model which is the
# neural network architecture built here.
def train_classifier(x_train, x_test, y_train, y_test, keras_model,
    ↪train_num=10):
    for count1 in range(0, 10):
        keras_model.fit(x_train, y_train, batch_size = 10, epochs=(train_num - 1),
    ↪verbose=0)
        keras_model.fit(x_train, y_train, batch_size = 10, epochs=1,
    ↪validation_data=(x_test, y_test), verbose=2)

```

```

[38]: # this function is used to make predictions on the test data set.
def predict_1(keras_model, x_test):
    y_pred = keras_model.predict(x_test)
    return y_pred

```

```

[39]: # Here we are preparing the input data and splitting it into train and test
x1, y1 = prepare_xy(
    df_PCA_output,
    [
        'residueCount', 'resolution', 'structureMolecularWeight',
        'crystallizationTempK', 'densityMatthews', 'densityPercentSol', 'phValue'
    ],
    ['PC1']
)
sc, x1s = scaler_1(x1)
x1s = pd.DataFrame(index=x1.index, columns=x1.columns, data=x1s)
x1_train, x1_test, y1_train, y1_test = splitter_1(x1s, y1, 0.1)

```

x: (65886, 7)

y: (65886, 1)

	PC1
-69696.566349	9
-40586.367926	7
-86536.188918	6
-86537.168902	5
-86554.198676	4
...	...
-79200.268356	1
-78995.971125	1
-83953.497492	1
-84714.907141	1

-93262.574191 1

[65705 rows x 1 columns]

x\_train: (59297, 7)

	residueCount	resolution	structureMolecularWeight	\
structureId				
1LL7	-0.016600	-0.265744	-0.045554	
1FV2	-0.207255	0.591600	-0.154751	
5GSU	0.285271	1.620412	0.346139	
4DFI	-0.417465	-0.608681	-0.300258	
3REH	0.277938	0.591600	0.339800	
...	...	...	...	
3L5H	-0.135759	2.477755	-0.101584	
4JVR	-0.319693	-0.780150	-0.227082	
2G75	0.064062	0.214368	-0.010429	
5UDV	0.552922	0.797362	0.315303	
409C	1.425536	-0.265744	0.778229	

	crystallizationTempK	densityMatthews	densityPercentSol	\
structureId				
1LL7	0.754795	-0.633587	-0.675803	
1FV2	0.206597	-0.165904	0.041208	
5GSU	0.206597	0.018334	0.753205	
4DFI	-0.012683	-1.129614	-1.699674	
3REH	-0.231962	-0.010010	0.238762	
...	...	...	...	
3L5H	0.754795	4.170789	2.661558	
4JVR	-1.547638	-0.704448	-0.809177	
2G75	0.206597	-0.463521	-0.393009	
5UDV	0.332682	-0.364315	-0.246599	
409C	0.425876	0.230917	0.517544	

	pHValue
structureId	
1LL7	0.552108
1FV2	1.327414
5GSU	-0.610851
4DFI	1.327414
3REH	-0.610851
...	...
3L5H	-0.610851
4JVR	-1.386157
2G75	1.327414
5UDV	-0.223198
409C	0.552108

[59297 rows x 7 columns]

x\_test: (6589, 7)

	residueCount	resolution	structureMolecularWeight	\
structureId				
1Y7A	0.053063	-0.660122	-0.019047	
4V9D	12.294720	1.448943	14.615002	
3EII	0.012732	0.162928	-0.026963	
2Q67	-0.356357	0.248662	-0.259335	
3MJ3	-0.469406	1.620412	-0.298657	
...	...	...	...	
1RK5	-0.192589	-0.608681	-0.162419	
3RL7	-0.063041	0.248662	-0.076680	
1Q7Q	0.196054	1.620412	0.092361	
1HZT	-0.379578	-1.208822	-0.273429	
4J1G	0.106226	1.088859	0.063445	

	crystallizationTempK	densityMatthews	densityPercentSol	\
structureId				
1Y7A	0.206597	0.755289	0.622840	
4V9D	-0.012683	0.882838	1.116223	
3EII	0.754795	-0.265110	-0.103197	
2Q67	0.754795	3.150391	2.335643	
3MJ3	0.206597	0.726944	0.982849	
...	...	...	...	
1RK5	0.425876	0.641911	0.884574	
3RL7	-0.451241	-0.831998	-1.054866	
1Q7Q	-1.547638	3.802312	2.552251	
1HZT	0.206597	-0.803653	-1.002719	
4J1G	-0.231962	-0.619415	-0.662766	

	pHValue
structureId	
1Y7A	2.102720
4V9D	-0.223198
3EII	-0.223198
2Q67	0.939761
3MJ3	0.552108
...	...
1RK5	-0.920974
3RL7	0.552108
1Q7Q	-0.843443
1HZT	-0.998504
4J1G	0.164455

[6589 rows x 7 columns]

y\_train: (59297, 1)

PC1

```

structureId
1LL7      -13158.819632
1FV2      -44703.197913
5GSU      99988.557954
4DFI      -86736.370547
3REH      98157.400246
...
3L5H      -29344.776423
4JVR      -65597.913477
2G75      -3011.869117
5UDV      91083.513759
409C      224812.153269

```

[59297 rows x 1 columns]

```

PC1
-69696.566349    9
-40586.367926    7
-86536.188918    6
-86537.168902    5
 241431.298559    4
...
-14268.240232    1
-382.630995      1
-63564.094675    1
-77513.224733    1
 224812.153269    1

```

[59144 rows x 1 columns]

y\_test: (6589, 1)

```

PC1
structureId
1Y7A      -5.501349e+03
4V9D      4.221809e+06
3EII      -7.788452e+03
2Q67      -7.491483e+04
3MJ3      -8.627429e+04
...
1RK5      -4.691825e+04
3RL7      -2.215034e+04
1Q7Q      2.668109e+04
1HZT      -7.898630e+04
4J1G      1.832764e+04

```

[6589 rows x 1 columns]

```

PC1
-66774.258365    2

```

```
-86554.198676      2
-5501.349088       1
 93348.899293       1
-83709.287793       1
...
-2890.129007        1
-44485.310474        1
-83900.786425        1
 132753.140476       1
 18327.640479        1
```

[6587 rows x 1 columns]

```
[40]: # Summary of the neural net architecture with the number of nodes
# scaled based on the number of input features
keras_model_1 = build_keras_model(x1s.shape[1], x1s.shape[1])
keras_model_1.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 4)	32
dense_1 (Dense)	(None, 1)	5
dense_2 (Dense)	(None, 4)	8
dense_3 (Dense)	(None, 7)	35

=====  
 Total params: 80  
 Trainable params: 80  
 Non-trainable params: 0  
 =====

```
[41]: # Training the model, This takes time.
# In the training the loss function value is an important paramater to follow
# as it gives an idea about how close the final representation is to
# the original data. The loss should decrease over multiple iterations (epochs)
train_classifier(
    x1_train, x1_test,
    x1_train, x1_test,
    keras_model_1, train_num=10,
)
```

5930/5930 - 14s - loss: 0.7948 - accuracy: 0.0000e+00 - binary\_accuracy:  
 0.0000e+00 - auc: 0.0000e+00 - true\_positives: 52065.0000 - true\_negatives:

0.0000e+00 - false\_positives: 0.0000e+00 - false\_negatives: 363014.0000 -  
 val\_loss: 0.7329 - val\_accuracy: 0.0000e+00 - val\_binary\_accuracy: 0.0000e+00 -  
 val\_auc: 0.0000e+00 - val\_true\_positives: 5946.0000 - val\_true\_negatives:  
 0.0000e+00 - val\_false\_positives: 0.0000e+00 - val\_false\_negatives: 40177.0000 -  
 14s/epoch - 2ms/step  
 5930/5930 - 15s - loss: 0.7923 - accuracy: 0.0000e+00 - binary\_accuracy:  
 0.0000e+00 - auc: 0.0000e+00 - true\_positives: 51471.0000 - true\_negatives:  
 0.0000e+00 - false\_positives: 0.0000e+00 - false\_negatives: 363608.0000 -  
 val\_loss: 0.7310 - val\_accuracy: 0.0000e+00 - val\_binary\_accuracy: 0.0000e+00 -  
 val\_auc: 0.0000e+00 - val\_true\_positives: 5763.0000 - val\_true\_negatives:  
 0.0000e+00 - val\_false\_positives: 0.0000e+00 - val\_false\_negatives: 40360.0000 -  
 15s/epoch - 3ms/step  
 5930/5930 - 14s - loss: 0.7911 - accuracy: 0.0000e+00 - binary\_accuracy:  
 0.0000e+00 - auc: 0.0000e+00 - true\_positives: 50301.0000 - true\_negatives:  
 0.0000e+00 - false\_positives: 0.0000e+00 - false\_negatives: 364778.0000 -  
 val\_loss: 0.7295 - val\_accuracy: 0.0000e+00 - val\_binary\_accuracy: 0.0000e+00 -  
 val\_auc: 0.0000e+00 - val\_true\_positives: 5870.0000 - val\_true\_negatives:  
 0.0000e+00 - val\_false\_positives: 0.0000e+00 - val\_false\_negatives: 40253.0000 -  
 14s/epoch - 2ms/step  
 5930/5930 - 14s - loss: 0.7905 - accuracy: 0.0000e+00 - binary\_accuracy:  
 0.0000e+00 - auc: 0.0000e+00 - true\_positives: 53594.0000 - true\_negatives:  
 0.0000e+00 - false\_positives: 0.0000e+00 - false\_negatives: 361485.0000 -  
 val\_loss: 0.7284 - val\_accuracy: 0.0000e+00 - val\_binary\_accuracy: 0.0000e+00 -  
 val\_auc: 0.0000e+00 - val\_true\_positives: 5874.0000 - val\_true\_negatives:  
 0.0000e+00 - val\_false\_positives: 0.0000e+00 - val\_false\_negatives: 40249.0000 -  
 14s/epoch - 2ms/step  
 5930/5930 - 14s - loss: 0.7904 - accuracy: 0.0000e+00 - binary\_accuracy:  
 0.0000e+00 - auc: 0.0000e+00 - true\_positives: 53667.0000 - true\_negatives:  
 0.0000e+00 - false\_positives: 0.0000e+00 - false\_negatives: 361412.0000 -  
 val\_loss: 0.7280 - val\_accuracy: 0.0000e+00 - val\_binary\_accuracy: 0.0000e+00 -  
 val\_auc: 0.0000e+00 - val\_true\_positives: 6186.0000 - val\_true\_negatives:  
 0.0000e+00 - val\_false\_positives: 0.0000e+00 - val\_false\_negatives: 39937.0000 -  
 14s/epoch - 2ms/step  
 5930/5930 - 14s - loss: 0.7902 - accuracy: 0.0000e+00 - binary\_accuracy:  
 0.0000e+00 - auc: 0.0000e+00 - true\_positives: 53915.0000 - true\_negatives:  
 0.0000e+00 - false\_positives: 0.0000e+00 - false\_negatives: 361164.0000 -  
 val\_loss: 0.7285 - val\_accuracy: 0.0000e+00 - val\_binary\_accuracy: 0.0000e+00 -  
 val\_auc: 0.0000e+00 - val\_true\_positives: 5853.0000 - val\_true\_negatives:  
 0.0000e+00 - val\_false\_positives: 0.0000e+00 - val\_false\_negatives: 40270.0000 -  
 14s/epoch - 2ms/step  
 5930/5930 - 14s - loss: 0.7900 - accuracy: 0.0000e+00 - binary\_accuracy:  
 0.0000e+00 - auc: 0.0000e+00 - true\_positives: 53333.0000 - true\_negatives:  
 0.0000e+00 - false\_positives: 0.0000e+00 - false\_negatives: 361746.0000 -  
 val\_loss: 0.7275 - val\_accuracy: 0.0000e+00 - val\_binary\_accuracy: 0.0000e+00 -  
 val\_auc: 0.0000e+00 - val\_true\_positives: 6107.0000 - val\_true\_negatives:  
 0.0000e+00 - val\_false\_positives: 0.0000e+00 - val\_false\_negatives: 40016.0000 -  
 14s/epoch - 2ms/step  
 5930/5930 - 15s - loss: 0.7908 - accuracy: 0.0000e+00 - binary\_accuracy:



```

0.0000e+00 - auc: 0.0000e+00 - true_positives: 52491.0000 - true_negatives:
0.0000e+00 - false_positives: 0.0000e+00 - false_negatives: 362588.0000 -
val_loss: 0.7289 - val_accuracy: 0.0000e+00 - val_binary_accuracy: 0.0000e+00 -
val_auc: 0.0000e+00 - val_true_positives: 5805.0000 - val_true_negatives:
0.0000e+00 - val_false_positives: 0.0000e+00 - val_false_negatives: 40318.0000 -
15s/epoch - 3ms/step
5930/5930 - 13s - loss: 0.7898 - accuracy: 0.0000e+00 - binary_accuracy:
0.0000e+00 - auc: 0.0000e+00 - true_positives: 54053.0000 - true_negatives:
0.0000e+00 - false_positives: 0.0000e+00 - false_negatives: 361026.0000 -
val_loss: 0.7275 - val_accuracy: 0.0000e+00 - val_binary_accuracy: 0.0000e+00 -
val_auc: 0.0000e+00 - val_true_positives: 6159.0000 - val_true_negatives:
0.0000e+00 - val_false_positives: 0.0000e+00 - val_false_negatives: 39964.0000 -
13s/epoch - 2ms/step
5930/5930 - 13s - loss: 0.7900 - accuracy: 0.0000e+00 - binary_accuracy:
0.0000e+00 - auc: 0.0000e+00 - true_positives: 53496.0000 - true_negatives:
0.0000e+00 - false_positives: 0.0000e+00 - false_negatives: 361583.0000 -
val_loss: 0.7274 - val_accuracy: 0.0000e+00 - val_binary_accuracy: 0.0000e+00 -
val_auc: 0.0000e+00 - val_true_positives: 6186.0000 - val_true_negatives:
0.0000e+00 - val_false_positives: 0.0000e+00 - val_false_negatives: 39937.0000 -
13s/epoch - 2ms/step

```

```

[42]: # After the model is trained we can make predictions. Here the predictions are
      ↪ on training set
      # meaning the neural network already learned this data and should be very good
      ↪ at predicting it again.
x1_train_pred = predict_1(keras_model_1, x1_train)
x1_train_pred = pd.DataFrame(index=x1_train.index, columns=x1_train.columns,
      ↪ data=x1_train_pred)
display(x1_train_pred)

```

```
1854/1854 [=====] - 3s 2ms/step
```

	residueCount	resolution	structureMolecularWeight \
structureId			
1LL7	0.000000e+00	0.000000e+00	0.000000e+00
1FV2	0.000000e+00	0.000000e+00	0.000000e+00
5GSU	1.604278e-02	9.999959e-01	4.211403e-03
4DFI	0.000000e+00	0.000000e+00	0.000000e+00
3REH	0.000000e+00	2.045232e-15	0.000000e+00
...	...	...	...
3L5H	6.441210e-08	9.997588e-01	3.130737e-10
4JVR	0.000000e+00	0.000000e+00	0.000000e+00
2G75	0.000000e+00	0.000000e+00	0.000000e+00
5UDV	0.000000e+00	0.000000e+00	0.000000e+00
409C	0.000000e+00	5.548220e-13	0.000000e+00

	crystallizationTempK	densityMatthews	densityPercentSol \
structureId			

1LL7	0.000000e+00	0.000000	0.000000
1FV2	0.000000e+00	0.000000	0.000000
5GSU	2.202764e-26	0.999997	0.999996
4DFI	0.000000e+00	0.000000	0.000000
3REH	0.000000e+00	0.040100	0.208307
...	...	...	...
3L5H	5.168652e-31	0.999988	0.999988
4JVR	0.000000e+00	0.000000	0.000000
2G75	0.000000e+00	0.000000	0.000000
5UDV	0.000000e+00	0.000000	0.000000
409C	0.000000e+00	0.223872	0.585437

	pHValue
structureId	
1LL7	0.000000e+00
1FV2	0.000000e+00
5GSU	5.847347e-01
4DFI	0.000000e+00
3REH	3.931682e-24
...	...
3L5H	1.164974e-02
4JVR	0.000000e+00
2G75	0.000000e+00
5UDV	0.000000e+00
409C	2.814206e-21

[59297 rows x 7 columns]

```
[43]: # prediction using test data subset.
x1_test_pred = predict_1(keras_model_1, x1_test)
x1_test_pred = pd.DataFrame(index=x1_test.index, columns=x1_test.columns,
                             data=x1_test_pred)
display(x1_test_pred)
```

206/206 [=====] - 1s 2ms/step

	residueCount	resolution	structureMolecularWeight \
structureId			
1Y7A	0.000000e+00	3.027744e-13	0.000000e+00
4V9D	9.999892e-01	9.999960e-01	9.997815e-01
3EII	0.000000e+00	0.000000e+00	0.000000e+00
2Q67	8.911152e-04	9.999894e-01	9.137367e-05
3MJ3	6.960497e-07	9.998894e-01	7.240248e-09
...	...	...	...
1RK5	0.000000e+00	7.112930e-11	0.000000e+00
3RL7	0.000000e+00	0.000000e+00	0.000000e+00
1Q7Q	1.230201e-05	9.999568e-01	3.205005e-07
1HZT	0.000000e+00	0.000000e+00	0.000000e+00

4J1G	0.000000e+00	0.000000e+00	0.000000e+00
------	--------------	--------------	--------------

	crystallizationTempK	densityMatthews	densityPercentSol
structureId			
1Y7A	0.000000e+00	0.189679	0.540786
4V9D	2.726076e-02	0.999985	0.999977
3EII	0.000000e+00	0.000000	0.000000
2Q67	1.826831e-27	0.999996	0.999995
3MJ3	3.972193e-30	0.999991	0.999990
...	...	...	...
1RK5	0.000000e+00	0.605994	0.858229
3RL7	0.000000e+00	0.000000	0.000000
1Q7Q	4.653240e-29	0.999994	0.999993
1HZT	0.000000e+00	0.000000	0.000000
4J1G	0.000000e+00	0.000000	0.000000

	phValue
structureId	
1Y7A	1.382845e-21
4V9D	2.406809e-01
3EII	0.000000e+00
2Q67	3.154294e-01
3MJ3	2.858829e-02
...	...
1RK5	8.360880e-19
3RL7	0.000000e+00
1Q7Q	8.153911e-02
1HZT	0.000000e+00
4J1G	0.000000e+00

[6589 rows x 7 columns]

```
[44]: # obtaining real values, or observables by using the final representation
# calculated by the neural network.
x1_train_inverse = pd.DataFrame(index=x1_train.index, columns=x1_train.columns,
    ↪data=sc.inverse_transform(x1_train))
x1_train_pred_inverse = pd.DataFrame(index=x1_train_pred.index,
    ↪columns=x1_train_pred.columns, data=sc.inverse_transform(x1_train_pred))
x1_test_inverse = pd.DataFrame(index=x1_test.index, columns=x1_test.columns,
    ↪data=sc.inverse_transform(x1_test))
x1_test_pred_inverse = pd.DataFrame(index=x1_test_pred.index,
    ↪columns=x1_test_pred.columns, data=sc.inverse_transform(x1_test_pred))

[45]: # displaying the newly calcualted features.
display(x1_train_inverse)
display(x1_train_pred_inverse)
display(x1_test_inverse)
```

```
display(x1_test_pred_inverse)
```

structureId	residueCount	resolution	structureMolecularWeight \
1LL7	784.0	2.00	87427.66
1FV2	472.0	2.50	55884.48
5GSU	1278.0	3.10	200574.00
4DFI	128.0	1.80	13852.53
3REH	1266.0	2.50	198742.88
...	...	...	...
3L5H	589.0	3.60	71242.50
4JVR	288.0	1.70	34990.44
2G75	916.0	2.28	97574.06
5UDV	1716.0	2.62	191666.55
409C	3144.0	2.00	325389.56

structureId	crystallizationTempK	densityMatthews	densityPercentSol	phValue
1LL7	298.00	2.23	44.85	7.5
1FV2	293.00	2.56	52.00	8.5
5GSU	293.00	2.69	59.10	6.0
4DFI	291.00	1.88	34.64	8.5
3REH	289.00	2.67	53.97	6.0
...	...	...	...	...
3L5H	298.00	5.62	78.13	6.0
4JVR	277.00	2.18	43.52	5.0
2G75	293.00	2.35	47.67	8.5
5UDV	294.15	2.42	49.13	6.5
409C	295.00	2.84	56.75	7.5

[59297 rows x 7 columns]

structureId	residueCount	resolution	structureMolecularWeight \
1LL7	811.164795	2.154981	100586.515625
1FV2	811.164795	2.154981	100586.515625
5GSU	837.418152	2.738175	101803.046875
4DFI	811.164795	2.154981	100586.515625
3REH	811.164795	2.154981	100586.515625
...	...	...	...
3L5H	811.164917	2.738037	100586.515625
4JVR	811.164795	2.154981	100586.515625
2G75	811.164795	2.154981	100586.515625
5UDV	811.164795	2.154981	100586.515625
409C	811.164795	2.154981	100586.515625

structureId	crystallizationTempK	densityMatthews	densityPercentSol \
-------------	----------------------	-----------------	---------------------

1LL7	291.115662	2.677063	51.589073
1FV2	291.115662	2.677063	51.589073
5GSU	291.115662	3.382668	61.560989
4DFI	291.115662	2.677063	51.589073
3REH	291.115662	2.705358	53.666302
...	...	...	...
3L5H	291.115662	3.382662	61.560902
4JVR	291.115662	2.677063	51.589073
2G75	291.115662	2.677063	51.589073
5UDV	291.115662	2.677063	51.589073
409C	291.115662	2.835029	57.427021

	phValue
structureId	
1LL7	6.787884
1FV2	6.787884
5GSU	7.542082
4DFI	6.787884
3REH	6.787884
...	...
3L5H	6.802910
4JVR	6.787884
2G75	6.787884
5UDV	6.787884
409C	6.787884

[59297 rows x 7 columns]

	residueCount	resolution	structureMolecularWeight \
structureId			
1Y7A	898.0	1.77	95084.64
4V9D	20931.0	3.00	4322348.00
3EII	832.0	2.25	92797.85
2Q67	228.0	2.30	25673.71
3MJ3	43.0	3.10	14315.06
...	...	...	...
1RK5	496.0	1.80	53669.27
3RL7	708.0	2.30	78436.41
1Q7Q	1132.0	3.10	127266.30
1HZZ	190.0	1.45	21602.38
4J1G	985.0	2.79	118913.50

	crystallizationTempK	densityMatthews	densityPercentSol	phValue
structureId				
1Y7A	293.0	3.21	57.80	9.5
4V9D	291.0	3.30	62.72	6.5
3EII	298.0	2.49	50.56	6.5
2Q67	298.0	4.90	74.88	8.0

3MJ3	293.0	3.19	61.39	7.5
...	...	...	...	...
1RK5	295.0	3.13	60.41	5.6
3RL7	287.0	2.09	41.07	7.5
1Q7Q	277.0	5.36	77.04	5.7
1HZT	293.0	2.11	41.59	5.5
4J1G	289.0	2.24	44.98	7.0

[6589 rows x 7 columns]

	residueCount	resolution	structureMolecularWeight	\
structureId				
1Y7A	811.164795	2.154981	100586.515625	
4V9D	2447.608398	2.738176	389388.375000	
3EII	811.164795	2.154981	100586.515625	
2Q67	812.623047	2.738172	100612.914062	
3MJ3	811.165894	2.738113	100586.515625	
...	...	...	...	
1RK5	811.164795	2.154981	100586.515625	
3RL7	811.164795	2.154981	100586.515625	
1Q7Q	811.184937	2.738153	100586.609375	
1HZT	811.164795	2.154981	100586.515625	
4J1G	811.164795	2.154981	100586.515625	

	crystallizationTempK	densityMatthews	densityPercentSol	\
structureId				
1Y7A	291.115662	2.810902	56.981762	
4V9D	291.364319	3.382659	61.560799	
3EII	291.115662	2.677063	51.589073	
2Q67	291.115662	3.382667	61.560978	
3MJ3	291.115662	3.382664	61.560928	
...	...	...	...	
1RK5	291.115662	3.104656	60.147289	
3RL7	291.115662	2.677063	51.589073	
1Q7Q	291.115662	3.382665	61.560951	
1HZT	291.115662	2.677063	51.589073	
4J1G	291.115662	2.677063	51.589073	

	pHValue
structureId	
1Y7A	6.787884
4V9D	7.098317
3EII	6.787884
2Q67	7.194729
3MJ3	6.824757
...	...
1RK5	6.787884
3RL7	6.787884

```

1Q7Q      6.893054
1HZT      6.787884
4J1G      6.787884

```

[6589 rows x 7 columns]

[46]: *# reindexing and relabeling for analyses in the next part*

```

x1_train_inverse.reset_index(inplace=True)
x1_train_inverse['label'] = 'train'
x1_train_pred_inverse.reset_index(inplace=True)
x1_train_pred_inverse['label'] = 'train_pred'
x1_test_inverse.reset_index(inplace=True)
x1_test_inverse['label'] = 'test'
x1_test_pred_inverse.reset_index(inplace=True)
x1_test_pred_inverse['label'] = 'test_pred'

```

[47]:

```

df_train_final = pd.concat([x1_train_inverse, x1_train_pred_inverse])
df_train_final.reset_index(inplace=True)
df_test_final = pd.concat([x1_test_inverse, x1_test_pred_inverse])
df_test_final.reset_index(inplace=True)

```

[48]:

```

display(df_train_final)
display(df_test_final)

```

	index	structureId	residueCount	resolution	structureMolecularWeight	\
0	0	1LL7	784.000000	2.000000	87427.660000	
1	1	1FV2	472.000000	2.500000	55884.480000	
2	2	5GSU	1278.000000	3.100000	200574.000000	
3	3	4DFI	128.000000	1.800000	13852.530000	
4	4	3REH	1266.000000	2.500000	198742.880000	
...	...	...	...	...	...	
118589	59292	3L5H	811.164917	2.738037	100586.515625	
118590	59293	4JVR	811.164795	2.154981	100586.515625	
118591	59294	2G75	811.164795	2.154981	100586.515625	
118592	59295	5UDV	811.164795	2.154981	100586.515625	
118593	59296	409C	811.164795	2.154981	100586.515625	

	crystallizationTempK	densityMatthews	densityPercentSol	phValue	\
0	298.000000	2.230000	44.850000	7.500000	
1	293.000000	2.560000	52.000000	8.500000	
2	293.000000	2.690000	59.100000	6.000000	
3	291.000000	1.880000	34.640000	8.500000	
4	289.000000	2.670000	53.970000	6.000000	
...	...	...	...	...	
118589	291.115662	3.382662	61.560902	6.802910	
118590	291.115662	2.677063	51.589073	6.787884	
118591	291.115662	2.677063	51.589073	6.787884	
118592	291.115662	2.677063	51.589073	6.787884	

118593                    291.115662                    2.835029                    57.427021    6.787884

                  label  
0                train  
1                train  
2                train  
3                train  
4                train

...                ...  
118589    train\_pred  
118590    train\_pred  
118591    train\_pred  
118592    train\_pred  
118593    train\_pred

[118594 rows x 10 columns]

	index	structureId	residueCount	resolution	structureMolecularWeight	\
0	0	1Y7A	898.000000	1.770000	9.508464e+04	
1	1	4V9D	20931.000000	3.000000	4.322348e+06	
2	2	3EII	832.000000	2.250000	9.279785e+04	
3	3	2Q67	228.000000	2.300000	2.567371e+04	
4	4	3MJ3	43.000000	3.100000	1.431506e+04	
...	...	...	...	...	...	
13173	6584	1RK5	811.164795	2.154981	1.005865e+05	
13174	6585	3RL7	811.164795	2.154981	1.005865e+05	
13175	6586	1Q7Q	811.184937	2.738153	1.005866e+05	
13176	6587	1HZT	811.164795	2.154981	1.005865e+05	
13177	6588	4J1G	811.164795	2.154981	1.005865e+05	

	crystallizationTempK	densityMatthews	densityPercentSol	phValue	\
0	293.000000	3.210000	57.800000	9.500000	
1	291.000000	3.300000	62.720000	6.500000	
2	298.000000	2.490000	50.560000	6.500000	
3	298.000000	4.900000	74.880000	8.000000	
4	293.000000	3.190000	61.390000	7.500000	
...	...	...	...	...	
13173	291.115662	3.104656	60.147289	6.787884	
13174	291.115662	2.677063	51.589073	6.787884	
13175	291.115662	3.382665	61.560951	6.893054	
13176	291.115662	2.677063	51.589073	6.787884	
13177	291.115662	2.677063	51.589073	6.787884	

                  label  
0                test  
1                test  
2                test  
3                test



```

4          test
...
13173 test_pred
13174 test_pred
13175 test_pred
13176 test_pred
13177 test_pred

```

```
[13178 rows x 10 columns]
```

After the training of the neural net is completed, I looked at the predictions of the network based on the training data subset and compared it to the original values. I do it by plotting their distributions and also comparing the means of the data (original vs prediction). The mean values show that the neural network can learn the hidden features of the data well enough to make predictions very close to original values. However, when the distributions of each feature is inspected we can recognize that neural network learned important features that are already closer to the mean values of the original data (peak of the distributions of blue curves), meaning the orange curve (distribution of predictions) have peak or peaks that are closer to the peaks of the original data. This is an interesting observation as the neural network show more than one peak for most of the features, implying that there might be two separate clusters of proteins in the data set.

```

[49]: for feature in [
        'residueCount', 'resolution', 'structureMolecularWeight',
        'crystallizationTempK', 'densityMatthews', 'densityPercentSol', 'phValue'
    ]:
        sns.displot(
            data=df_train_final,
            x=feature,
            hue='label',
            kind='kde',
            common_norm=False,
            aspect=2,
        )
        print(
            '{} means for training data, original:{:.2f} <-> prediction:{:.2f}'.
            format(
                feature,
                df_train_final[df_train_final['label'] == 'train'][feature].mean(),
                df_train_final[df_train_final['label'] == 'train_pred'][feature].
                mean()
            )
        )

```

```

residueCount means for training data, original:812.90 <-> prediction:896.42
resolution means for training data, original:2.16 <-> prediction:2.25
structureMolecularWeight means for training data, original:100858.02 <->
prediction:112630.61
crystallizationTempK means for training data, original:291.09 <->

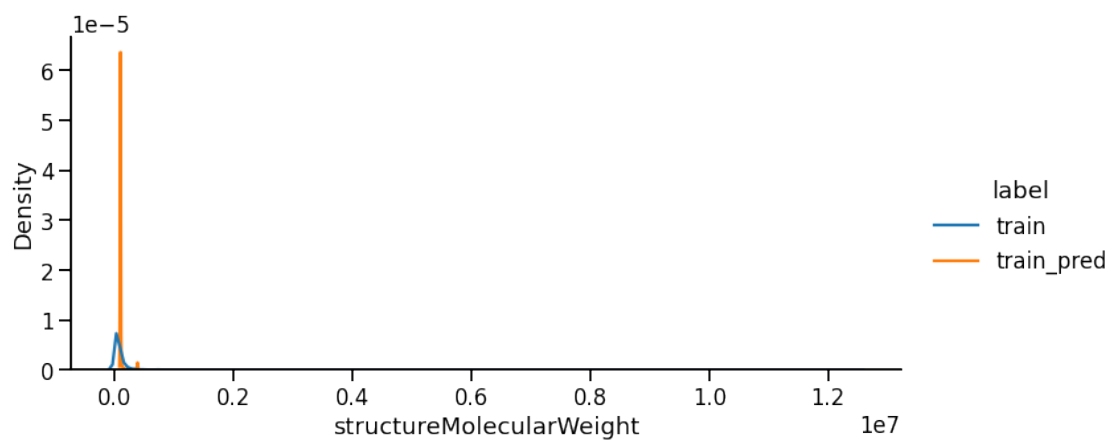


```


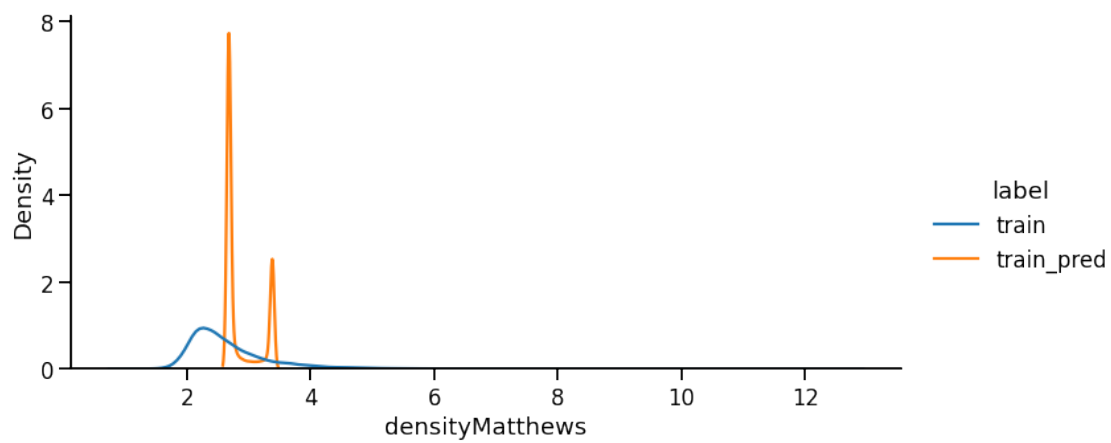

prediction:291.12


densityMatthews means for training data, original:2.68 <-> prediction:2.87

densityPercentSol means for training data, original:51.58 <-> prediction:54.79

phValue means for training data, original:6.79 <-> prediction:6.87




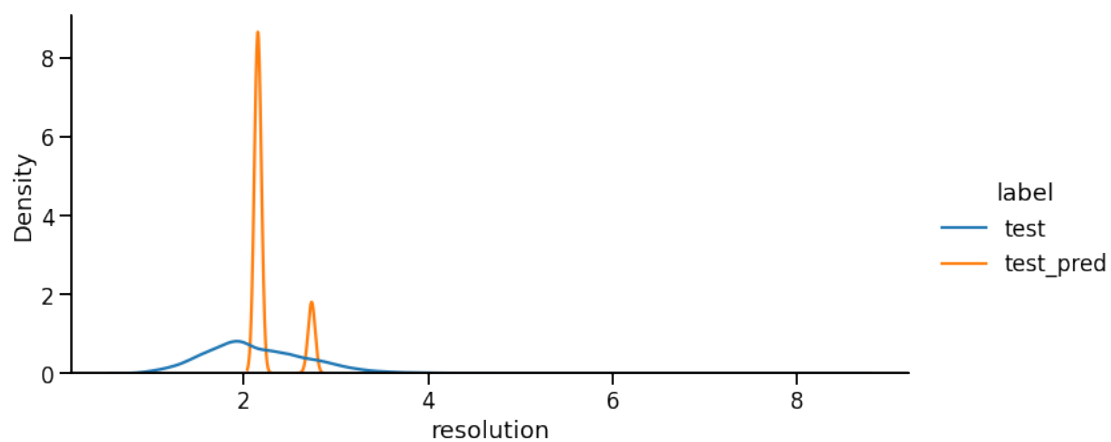



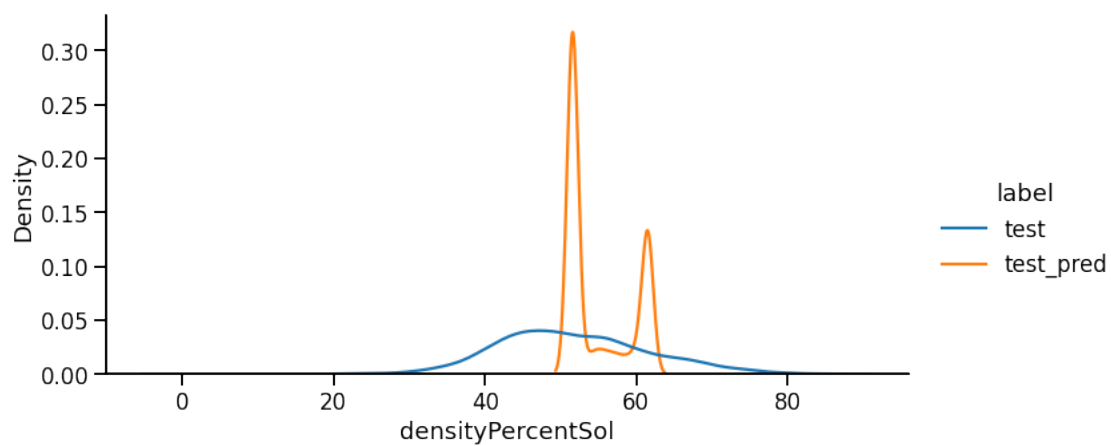
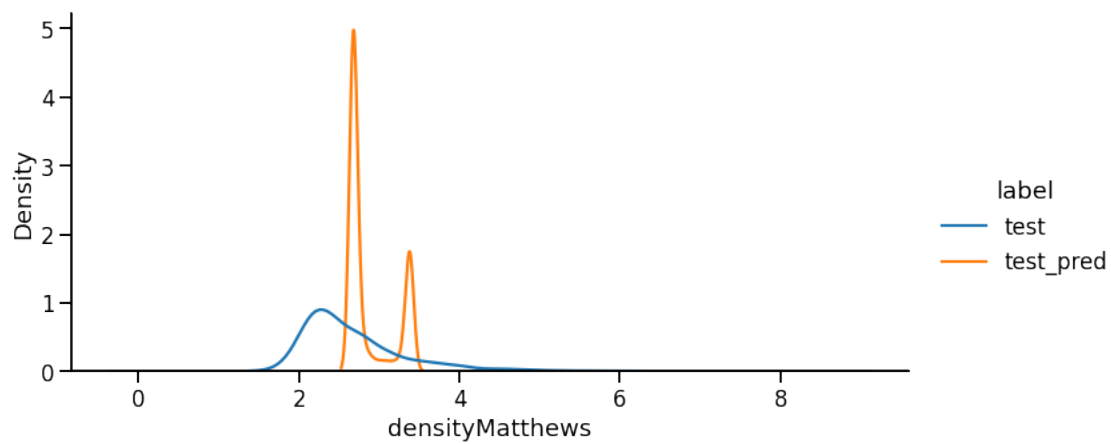
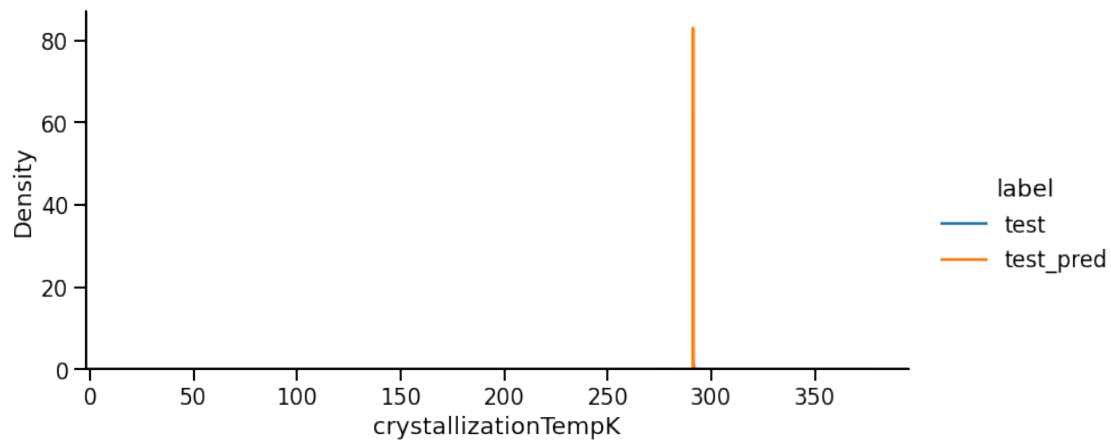



The same analyses applied on training data is subjected on the testing subset. This is very crucial as this is data which is completely new to the neural network (unseen) data. The results show that, we still observe a similar trend seen with the training set as the predictions are close to the original values within ~5% deviation.

```
[50]: for feature in [
        'residueCount', 'resolution', 'structureMolecularWeight',
        'crystallizationTempK', 'densityMatthews', 'densityPercentSol', 'pHValue'
    ]:
        sns.displot(
            data=df_test_final,
            x=feature,
            hue='label',
            kind='kde',
            common_norm=False,
            aspect=2,
        )
        print(
            '{} means for test data, original:{:.2f} <-> prediction:{:.2f}'.format(
                feature,
                df_test_final[df_test_final['label'] == 'test'][feature].mean(),
                df_test_final[df_test_final['label'] == 'test_pred'][feature].mean()
            )
        )
```

```
residueCount means for test data, original:795.59 <-> prediction:897.36
resolution means for test data, original:2.15 <-> prediction:2.26
structureMolecularWeight means for test data, original:98143.13 <->
prediction:112591.51
crystallizationTempK means for test data, original:291.31 <-> prediction:291.12
densityMatthews means for test data, original:2.68 <-> prediction:2.88
densityPercentSol means for test data, original:51.65 <-> prediction:54.89
pHValue means for test data, original:6.80 <-> prediction:6.87
```







Overall the neural network treatment was successful and it captured the essence of the structural protein features strong enough to make predictions not so far from the originals. However one drawback seen from distributions is that predictions are not as diverse as the original data, meaning predictions are mostly centered on one or two regions whereas original data points are distributed across multiple values for each feature. Besides the drawback neural network was able to identify multiple classes which can not be easily comprehended from the distribution of original values. This is the strongest achievement of the autoencoder as it captures the fundamental details that are missed by human (by human observation). On the other hand, when the computational cost is considered the overall enhancement in classification compared to PCA is minimal. While for the current data set this might be true but, considering only 7 features in structural protein data, when the number of features are increased several fold, the power of autoencoders will shine more as the first two components of PCA could become very convoluted and several more components should be considered and additional techniques should be incorporated to further understand the data.

## 7 Conclusion

Through my analysis of structural discovery protein data, we were able to utilize both machine learning, PCA and neural network techniques to gain new insights into several properties of the proteins. This study highlights the challenges of achieving results that resemble natural phenomena by comparing solution densities to crystal densities. By using machine learning techniques (k-means clustering, hierarchical clustering, DBSCAN) these results obtained show that only 26% of the experimental data points display naturally occurring behavior that captures the correct underlying biological fundamentals. Thus, I used PCA technique to better understand the data. PCA allowed me to identify key factors driving variation in the dataset and provided a simplified view of the underlying patterns and relationships. PCA analyses showed that the first three principal components can distinguish key protein features from each other and enabled me to cluster the protein labels based on these differences.

In addition, by using a non-linear approach like neural network, I was able to accurately distinguish input features, highlighting its potential for learning complex patterns in protein structural discovery data. The results showed that the autoencoder learned the hidden characteristics of protein structural features and was able to reconstruct the features accurately. When compared

to the distributions and the means of the original data, the features that were reconstructed from the autoencoder showed high overlap and indicated the existence of two separate clusters (multiple peaks in the distribution) of proteins in the dataset. This observation is intriguing as for future investigations to determine if there are indeed distinct groups of proteins provided their structural features. Overall, the results demonstrate the potential of both PCA and neural networks in analyzing complex protein data, and the insights gained from these techniques can provide valuable information for further research in the field.

## 8 References

- “Structural Protein Sequences.” n.d. <https://www.kaggle.com/datasets/shahir/protein-dataset>.
- Warren, Bertram Eugene. 1990. X-Ray Diffraction. Courier Corporation.
- Whittig, L. D., and W. R. Allardice. 1986. “X-Ray Diffraction Techniques.” *Methods of Soil Analysis: Part 1 Physical and Mineralogical Methods* 5: 331362.
- Fausett, L. V. (2006). *Fundamentals of neural networks: architectures, algorithms and applications*. Pearson Education India.
- Haykin, S. (2009). *Neural networks and learning machines*, 3/E. Pearson Education India.
- SKLEARN.DECOMPOSITION.PCA. scikit. (n.d.). Retrieved February 28, 2023, from <https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.PCA.html>